

UCS503 — Software Engineering Project
Academic Year 2026–2027

Solar Sage AI

Drone-Based Edge-AI System for Condition-Based Solar Panel Cleaning

Dynamic Perception, Edge Decision-Making and Targeted Actuation

Mid-Semester Evaluation Report

Submitted by

1024030752 Shorya Gupta

1024030757 Omkar Shukla

1024030760 Sahil Soni

B.E. Third Year, Computer Science & Engineering

Submitted to

Dr. Jeelani Asif

Department of Computer Science and Engineering

Thapar Institute of Engineering and Technology

Patiala, Punjab, India

Table of Contents

Chapter 1 Introduction	3
1.1 Background and Context	3
1.1.1 Problem Statement	3
1.2 Project Objectives	3
Chapter 2 Methodology and Design	4
2.1 System Architecture	4
2.1.1 High-Level Design	4
2.1.2 Technical Stack	4
2.2 Implementation Details	5
2.2.1 Module 1 — Core Perception & Dirt Detection Pipeline	5
2.2.2 Module 2 — Edge Decision Layer & Actuation Control	5
Chapter 3 Results and Evaluation	6
3.1 Testing Methodology	6
3.2 Performance Metrics	6
3.3 Analysis	7
Chapter 4 Conclusion	7
4.1 Summary of Work	7
4.2 Key Findings	7
4.3 Future Work and Recommendations	7
4.4 Final Remarks	7
Bibliography	8

Chapter 1 Introduction

1.1 Background and Context

Solar operators at utility scale routinely face a mismatch between how panels are maintained and how dirty they actually are. Fixed-schedule manual cleaning consumes water and labour, can expose panels to repeated physical contact, and does not prioritize cleaning according to the actual condition of individual panels.

Solar Sage AI addresses this operational gap through a drone-based “Inspect-to-Clean” system. An RGB/thermal drone payload captures panel imagery, an OpenCV-based perception layer isolates panels and dirt or debris regions, and an edge decision layer converts the observed condition into a panel-level dirt score and cleaning priority. The execution layer then uses ESP32-controlled servo and pump signals to perform targeted cleaning.

1.1.1 Problem Statement

- **Resource intensity:** manual cleaning at scale consumes substantial water and labour; the project proposal cites a 1 GW farm using upwards of 12.5 million litres per month on washing alone.
- **Schedule blindness:** fixed-calendar cleaning can service panels that are still reasonably clean while genuinely soiled panels remain unaddressed.
- **Physical degradation risk:** frequent contact-based cleaning can contribute to micro-cracks and long-term efficiency loss.
- **Fragmented decision-making:** without per-panel condition data, operators cannot prioritize cleaning according to actual impact.

1.2 Project Objectives

- **Dynamic perception:** capture RGB/thermal imagery and segment panel surfaces using OpenCV-based detection.
- **Condition-based scoring:** assign each panel a simple 0–100 dirt score grouped into Low, Medium and High severity.
- **Targeted actuation:** issue precise servo and pump commands only when panels cross the configured dirt threshold.
- **Edge-first response:** maintain a perception-to-decision latency target below one second with on-edge processing.
- **Operational measurement:** quantify water use, cleaning time, output gain, cost savings and ROI.

Chapter 2 Methodology and Design

2.1 System Architecture

2.1.1 High-Level Design

Solar Sage AI utilises a modular, three-tier architecture: the Client Dashboard Layer, the Core Perception & Decision Engine, and the Actuation & Telemetry Store. Figure 2.1 illustrates the high-level architecture.



Figure 2.1 — Solar Sage AI High-Level Multi-Tier Architecture

2.1.2 Technical Stack

- **Backend Runtime & API:** Python 3.11, Flask REST API with a Redis/Celery job queue for flight and processing orchestration.
- **Perception & Computer Vision:** OpenCV (colour/texture thresholding), NumPy and Pillow for RGB/thermal frame processing and panel segmentation.
- **Edge Firmware:** ESP32 (Arduino/ESP-IDF), PWM servo control, and a relay-driven water-pump actuation circuit.
- **Decision Layer:** Python heuristic scoring engine implementing confidence, priority, benchmarking, ROI, and water-optimisation agents.
- **Database & Storage:** SQLite for the prototype (PostgreSQL planned for production), with JSON/CSV telemetry logs for latency and trigger-accuracy measurement.
- **Frontend UI:** React, Tailwind CSS, and Chart.js for per-panel dirt-score and ROI visualisation.

2.2 Implementation Details

2.2.1 Module 1 — Core Perception & Dirt Detection Pipeline

The perception module flies a fixed inspection pattern over the test array, captures RGB/thermal frames, and segments each frame to isolate individual panels before flagging dirt or debris regions using colour and texture feature thresholds (Figure 2.2).




Figure 2.2 — Core Perception & Dirt Detection Pipeline

2.2.2 Module 2 — Edge Decision Layer & Actuation Control

The decision layer converts flagged dirt regions into a 0–100 dirt score per panel, groups panels into Low/Medium/High priority, and suppresses low-confidence detections before any actuation signal is issued. Confirmed commands are passed to the ESP32 execution layer, which drives servo positioning and pump timing for targeted water application (Figure 2.3).



Figure 2.3 — Edge Decision Layer & Actuation Control Flow

The multi-agent decision layer retains the five agents specified in the Solar Sage AI prototype: Confidence Agent, Priority Agent, Benchmarking Agent, Impact & ROI Agent, and Water & Action Optimizer.

Chapter 3 Results and Evaluation

3.1 Testing Methodology

Testing followed a structured, multi-tier evaluation framework:

- **Unit Testing:** individual testing of OpenCV thresholding routines, the dirt-score calculator, and ESP32 servo/pump command encoding.
- **Integration Testing:** end-to-end testing of the detect-and-clean loop under simulated lighting and glare conditions.
- **Benchmark Validation:** validation across realistic soiling scenarios on the test array — Low soiling, Medium soiling, and High soiling.




Figure 3.1 — Multi-Tier Testing and Quality Assurance Framework

3.2 Performance Metrics

The prototype underwent quantitative benchmarking to measure latency, trigger accuracy, water use, cleaning time, system uptime, and concurrent processing throughput. Table 3.1 summarizes the target versus achieved engineering metrics.

Evaluation Metric	Target Specification	Achieved (Prototype v0.3)	Status
Perception-to-Decision Latency	≤ 1.0 s	0.71 s	Exceeded
Cleaning Trigger Accuracy	≥ 90.0%	93.8%	Exceeded
Water Use per Panel (vs. 2.8 L manual baseline)	≤ 0.6 L	0.34 L	Exceeded
Cleaning Time per Panel	≤ 15 s	5.2 s	Exceeded
Dashboard / Control-Loop Uptime	≥ 99.0%	99.4%	Exceeded
Concurrent Processing Throughput	≥ 20 panels/min	34 panels/min	Exceeded

Table 3.1 — Solar Sage AI Prototype Performance Metrics Summary

3.3 Analysis

The prototype results indicate that the Solar Sage AI workflow can satisfy the core engineering targets defined for the pilot. Perception-to-decision latency of 0.71 s remains below the one-second target, while cleaning trigger accuracy reaches 93.8%. The reported water-use figure of 0.34 L per panel is substantially below the 2.8 L manual baseline and the 0.6 L project target.

The prototype also reports a 5.2-second cleaning time per panel, 99.4% dashboard/control-loop uptime, and concurrent processing of 34 panels per minute. Together, these figures support the intended condition-based workflow: inspect the array, identify panels that require attention, and actuate cleaning only where the decision layer indicates a sufficient need.

Chapter 4 Conclusion

4.1 Summary of Work

During the first half of the project lifecycle, the team defined and prototyped an end-to-end Solar Sage AI workflow connecting drone inspection, RGB/thermal image capture, OpenCV perception, heuristic dirt scoring, multi-agent decision-making, ESP32 actuation, and telemetry/dashboard reporting.

4.2 Key Findings

- **Fast edge response:** 0.71 s perception-to-decision latency was achieved in the prototype benchmark.
- **Reliable cleaning decisions:** cleaning trigger accuracy reached 93.8%, above the project target.
- **Reduced water use:** the Prototype v0.3 figure is 0.34 L per panel versus a 2.8 L manual baseline.
- **Targeted actuation:** 5.2 s cleaning time per panel supports short, condition-based intervention.
- **Operational readiness:** 99.4% control-loop/dashboard uptime and 34 panels/min processing throughput were reported.

4.3 Future Work and Recommendations

- Extend the multi-agent decision layer for confidence, priority, benchmarking, ROI, and water optimisation across larger arrays.
- Validate the system across larger test arrays and wider ranges of dust, lighting, glare, weather, and panel orientation.
- Move from heuristic dirt classification toward learned edge-AI models as a stretch goal for difficult visual cases.
- Strengthen multi-drone coordination and scalable dashboard data handling for deployment at utility scale.
- Containerize the software stack and formalize CI/CD staging and production deployment.

4.4 Final Remarks

Solar Sage AI offers a practical engineering path from inspection to condition-based cleaning by linking perception, decision-making, actuation, and measurement in one control loop. The prototype demonstrates measurable improvements in latency, trigger accuracy, resource consumption, and operational throughput, while the architecture leaves a clear path toward larger pilot deployments and more advanced edge-AI perception.